

K-SVD Dictionary Learning: Reproduction and Extension to Image Denoising

ECE 417 / ECE 613 Final Project — Winter 2026

Yifan Zhou

Department of Electrical and Computer Engineering
University of Waterloo, Waterloo, ON, Canada

Abstract—We reproduce the K-SVD algorithm of Aharon, Elad, and Bruckstein [1], which learns an over-complete dictionary from training signals using an alternating sparse-coding and rank-1 SVD update strategy. Following the paper’s synthetic dictionary-recovery experiment, our implementation recovers 45.5 ± 2.2 out of 50 atoms in the noiseless case (vs. 44.1 ± 1.9 for the MOD baseline), and performs comparably or better than MOD across all tested noise levels, reproducing the paper’s key qualitative finding. As a novel extension, we apply the learned-dictionary framework to image denoising—a problem not addressed in the original paper—and compare against Non-Local Means (NLM). We further propose a combined NLM+K-SVD pipeline where NLM pre-filtering provides cleaner training data for the dictionary, which improves over standalone K-SVD on smooth images and consistently outperforms NLM across both test images.

Index Terms—Dictionary learning, K-SVD, sparse representation, orthogonal matching pursuit, image denoising, overcomplete dictionaries.

I. INTRODUCTION AND BACKGROUND

Sparse representation has become a foundational tool in signal processing and image analysis. The central idea is that natural signals can be well approximated as sparse linear combinations of a small set of *atoms* drawn from an over-complete *dictionary* [4]. Early work used fixed, analytically designed dictionaries such as wavelets and DCT bases, which are efficient but not adapted to the signal class of interest.

Dictionary learning addresses this limitation by jointly optimising the dictionary and the sparse coefficients from a set of training signals. Olshausen and Field [5] showed that learning a dictionary from image patches recovers Gabor-like features resembling receptive fields in primary visual cortex. The Method of Optimal Directions (MOD) [2] formulated dictionary learning as an alternating minimisation, updating the dictionary globally via the pseudoinverse.

Aharon, Elad, and Bruckstein [1] introduced K-SVD, a more efficient alternating algorithm that updates each dictionary atom individually using a rank-1 singular value decomposition (SVD) of the corresponding residual matrix. K-SVD generalises K-means clustering and has been applied to image compression, inpainting, and—in follow-up work [6]—image denoising. For denoising, Non-Local Means (NLM) [8] represents a classical baseline that exploits non-local self-similarity rather than a learned dictionary.

This project reproduces the synthetic dictionary-recovery experiment of [1] (Section V-A), comparing K-SVD with MOD as the paper’s primary baseline. We then extend K-SVD to image denoising—a problem the original paper does not address—compare it against NLM, and propose a combined NLM+K-SVD pipeline that improves over both.

II. REPRODUCIBLE COMPONENT: K-SVD ALGORITHM

A. Algorithm

The K-SVD algorithm [1] solves:

$$\min_{\mathbf{D}, \mathbf{X}} \|\mathbf{Y} - \mathbf{DX}\|_F^2 \quad \text{s.t.} \quad \forall i: \|\mathbf{x}_i\|_0 \leq L, \quad (1)$$

where $\mathbf{Y} \in \mathbb{R}^{n \times N}$ is the training matrix, $\mathbf{D} \in \mathbb{R}^{n \times K}$ the dictionary, and $\mathbf{X} \in \mathbb{R}^{K \times N}$ the sparse codes. K-SVD alternates between two steps.

Step 1 – Sparse Coding. With $\mathbf{D}$ fixed, each $\mathbf{y}_i$ is encoded by OMP [3] with sparsity L .

Step 2 – Dictionary Update. For each atom $\mathbf{d}_k$, let $\Omega_k = \{i: x_{k,i} \neq 0\}$ and form the partial residual

$$\mathbf{E}_k = \mathbf{Y}_{\Omega_k} - \sum_{j \neq k} \mathbf{d}_j \mathbf{x}_j^\top|_{\Omega_k}. \quad (2)$$

The SVD $\mathbf{E}_k = \mathbf{U}\Sigma\mathbf{V}^\top$ gives the updated atom $\mathbf{d}_k \leftarrow \mathbf{u}_1$ and updated coefficients $\mathbf{x}_k|_{\Omega_k} \leftarrow \sigma_1 \mathbf{v}_1$. This rank-1 step guarantees non-increasing reconstruction error.

MOD [2] uses the same OMP step but replaces the column-by-column update with a single global pseudoinverse step: $\mathbf{D} \leftarrow \mathbf{Y}\mathbf{X}^\dagger$.

B. Experimental Setup

We follow Section V-A of [1] exactly: a true dictionary $\mathbf{D}_0 \in \mathbb{R}^{20 \times 50}$ with unit-norm Gaussian columns; $N = 1500$ signals each built from $L = 3$ randomly chosen atoms with Gaussian weights; additive Gaussian noise at $\text{SNR} \in \{\infty, 30, 20, 10\}$ dB; 80 iterations; 20 Monte Carlo trials. An atom $\hat{\mathbf{d}}$ is *recovered* when $\max_k |\hat{\mathbf{d}}^\top \mathbf{d}_{0,k}| \geq 1 - \varepsilon$ ($\varepsilon = 0.01$).

C. Results and Discussion

Table I and Fig. 1 show that K-SVD matches or exceeds MOD at all noise levels, consistent with the paper’s claim. Our absolute counts (44–46) are 3–5 atoms below the paper’s reported values (~ 28 –49). We attribute this to three factors:

TABLE I
ATOMS RECOVERED OUT OF 50 (MEAN $\pm$ STD, 20 TRIALS).

Method	No noise	30 dB	20 dB	10 dB
K-SVD	45.5 $\pm$ 2.2	44.6 $\pm$ 2.2	44.8 $\pm$ 1.9	45.1 $\pm$ 3.0
MOD	44.1 $\pm$ 1.9	45.3 $\pm$ 1.6	44.5 $\pm$ 2.2	44.8 $\pm$ 2.4
Paper (K-SVD) [1]	$\sim$ 49	$\sim$ 46	$\sim$ 40	$\sim$ 28




Fig. 1. Atoms recovered by K-SVD and MOD at four noise levels. Error bars: ± 1 std over 20 trials.

(1) fewer Monte Carlo trials (20 vs. 50 in the paper), increasing variance; (2) the dictionary update averages over ~ 90 activating signals per atom, suppressing additive noise and reducing sensitivity to SNR; and (3) OMP conditioning issues during early training iterations, which the paper’s MATLAB solver may handle more robustly. Critically, the paper’s central qualitative finding—K-SVD outperforms MOD and remains useful under noise—is reproduced.

D. Effect of Training Set Size

Fig. 2 extends the evaluation to vary the number of training signals $N \in \{100, 300, 500, 1000, 1500, 2000, 3000\}$ in the noiseless setting. Recovery rises sharply from near zero at $N = 100$ to ~ 40 atoms at $N = 500$, then levels off near 45–46 above $N = 1500$, matching the trend of paper Figure 6. A notable crossover occurs around $N = 500$: MOD leads K-SVD at $N = 300$ (26.5 vs. 18.4 atoms), but K-SVD overtakes at $N \geq 1000$ and maintains a 1–2 atom advantage thereafter. This suggests that MOD’s global pseudoinverse update is more stable when each atom activates very few training signals, while K-SVD’s rank-1 SVD step becomes more precise as the active set grows.

E. Convergence Analysis

We further track the per-iteration reconstruction error $\|\mathbf{Y} - \mathbf{DX}\|_F^2/N$ for both methods over 80 iterations (5 trials, noiseless). Fig. 3 reveals a nuanced finding: MOD actually converges to a *lower* reconstruction error (~ 0.040) than K-SVD (~ 0.044) within 80 iterations, while K-SVD continues to decrease slowly.

This seeming contradiction with K-SVD’s superior atom recovery is instructive: *reconstruction error is not the same as dictionary quality*. MOD’s pseudoinverse step is the globally




Fig. 2. Atoms recovered vs. number of training signals N (noiseless; other settings as in Table I). Error bars: ± 1 std over 20 trials.




Fig. 3. Reconstruction error vs. iteration for K-SVD and MOD. Mean over 5 trials.

optimal dictionary update given fixed sparse codes, which aggressively minimises reconstruction error but may converge to a dictionary that represents the training data well without structurally matching the true atoms. K-SVD’s column-by-column SVD update is less aggressive on reconstruction error but preserves a tighter correspondence to the underlying atoms, yielding higher atom recovery (Table I). This trade-off clarifies why the two metrics can diverge and supports the paper’s choice of atom recovery as the primary evaluation criterion.

III. NOVELTY COMPONENT: K-SVD FOR IMAGE DENOISING

A. Motivation

The original paper [1] demonstrates K-SVD on synthetic signals, facial compression, and a brief inpainting demo, but does *not* address image denoising. We extend K-SVD to this problem—a different signal recovery task—to assess how well the learned dictionary framework generalises beyond the paper’s studied applications.

B. Method

Given a noisy image $\mathbf{v} = \mathbf{x} + \boldsymbol{\eta}$ ($\boldsymbol{\eta} \sim \mathcal{N}(0, \sigma^2 \mathbf{I})$), we extract all overlapping 8×8 patches (stride 1), mean-subtract each patch, and train a 64×256 dictionary via mini-batch K-SVD [7] on 50 000 randomly sampled patches. Each patch

is then re-encoded with OMP under a noise-adapted stopping criterion [6]:

$$\hat{\alpha} = \arg \min_{\alpha} \|\alpha\|_0 \text{ s.t. } \|\mathbf{D}\alpha - \mathbf{y}\|_2^2 \leq C\sigma^2 n, \quad (3)$$

with $C = 1.15$ and $n = 64$. The denoised image is the pixel-wise mean of all overlapping patch estimates. We compare against NLM [8].

C. Results

TABLE II
DENOISING PSNR (dB) ON CAMERA MAN (512×512). BEST PER ROW IN **BOLD**.

σ	Noisy	K-SVD	NLM
10	28.24	33.31	32.86
20	22.41	29.77	29.53
25	20.59	28.75	28.71
30	19.14	27.87	28.04




Fig. 4. Visual comparison at $\sigma = 25$. Left to right: original, noisy (20.59 dB), K-SVD (28.75 dB), NLM (28.71 dB).




Fig. 5. The 256 dictionary atoms learned by K-SVD from Camera Man patches (8×8 , $K = 256$, $\sigma = 25$). Atoms resemble oriented Gabor filters and step edges, confirming that the learned dictionary adapts to natural image structure.

TABLE III
DENOISING PSNR (dB) ON ASTRONAUT (512×512 , GRAYSCALE). BEST PER ROW IN **BOLD**.

σ	Noisy	K-SVD	NLM
10	28.53	34.64	33.40
20	22.61	30.40	29.64
25	20.75	28.95	28.34
30	19.25	27.72	27.22

D. Generalisation Across Images

Table III repeats the comparison on the Astronaut test image (512×512 , grayscale), which differs in content and texture from Camera Man. K-SVD outperforms NLM by 0.5–1.2 dB across all noise levels, a larger gap than on Camera Man. The relative ordering $\text{K-SVD} > \text{NLM}$ is preserved across both images, confirming that the findings generalise beyond a single test case. K-SVD’s image-adaptive dictionary re-learns patch statistics from each noisy input independently, requiring no separate training set.

E. Ablation: Dictionary Size

TABLE IV
PSNR VS. DICTIONARY SIZE K AT $\sigma = 25$.

K	PSNR (dB)
64	28.59
128	28.65
256	28.75
512	28.80

Doubling K from 64 to 512 yields only +0.21 dB (Table IV), indicating representational saturation for 8×8 natural-image patches near $K = 256$. OMP cost scales linearly with K , so larger dictionaries offer an unfavourable compute-quality trade-off.

F. NLM-Guided Dictionary Learning

A key limitation of standalone K-SVD is that its dictionary is trained on noisy patches, limiting atom quality. We propose a two-step extension: (1) apply NLM to obtain a coarse estimate, then (2) train the K-SVD dictionary on the cleaner NLM patches while encoding the *original noisy* patches with the standard OMP threshold. This preserves the correct noise model for encoding while supplying cleaner training data for dictionary learning.



Fig. 6. Visual comparison of the proposed NLM+K-SVD against baselines on Camera Man at $\sigma = 25$.

Table V shows that NLM+K-SVD outperforms *both* standalone K-SVD and NLM on the Astronaut image at *all* four

TABLE V
PSNR (dB): NLM+K-SVD vs. K-SVD and NLM. BEST PER ROW IN
BOLD.

Image	σ	Noisy	NLM+K-SVD	K-SVD	NLM
Cam.Man	10	28.24	33.32	33.31	32.86
	20	22.41	29.77	29.75	29.53
	25	20.59	28.70	28.75	28.71
	30	19.14	27.75	27.87	28.04
Astronaut	10	28.53	34.66	34.64	33.40
	20	22.61	30.42	30.40	29.64
	25	20.75	28.99	28.95	28.34
	30	19.25	27.78	27.72	27.22

noise levels, with gains of 0.02–0.06 dB over K-SVD and 0.56–1.26 dB over NLM. On Camera Man the gain is marginal at $\sigma = 10$ and slightly negative at higher noise, suggesting that the benefit depends on image content: NLM is more effective on smooth regions (Astronaut) than on high-frequency textures (Camera Man), leading to a cleaner training set and a higher-quality dictionary. The consistent improvement on Astronaut confirms that combining non-local pre-filtering with dictionary learning is a viable direction for further development.

G. Analysis

K-SVD matches or slightly exceeds NLM at $\sigma \leq 25$, where the learned dictionary captures fine image-specific structure that NLM’s fixed patch-matching misses. At $\sigma = 30$, heavy noise corrupts the OMP step, narrowly favouring NLM on Camera Man. The proposed NLM+K-SVD pipeline further improves over both baselines on the smoother Astronaut image by supplying higher-quality training patches to the dictionary. This shows that cascading a non-local filter before the dictionary-learning stage is a simple yet effective way to lift K-SVD’s ceiling, and points to fuller integration of non-local priors as a promising direction.

IV. DISCUSSION AND CONCLUSIONS

Reproducibility. Our implementation faithfully reproduces the K-SVD algorithm and its primary finding: column-by-column SVD updates (K-SVD) yield better dictionary quality than global pseudoinverse updates (MOD), particularly in noisy settings. The 3–5 atom gap relative to the paper is explained by fewer Monte Carlo trials, the noise-suppressing effect of averaging over many activating signals, and OMP solver differences; it does not alter the qualitative conclusion.

Generalisation to denoising. Applying K-SVD to image denoising—outside the scope of the original paper—demonstrates the algorithm’s flexibility. Competitive or superior performance relative to NLM validates the utility of adaptive dictionary learning for image restoration. The proposed NLM+K-SVD pipeline further improves over both individual components on smooth images (Astronaut: +0.02–0.06 dB over K-SVD and +0.56–1.26 dB over NLM at all noise levels), confirming that feeding cleaner training patches to the dictionary is beneficial.

Future directions. Extending the NLM+K-SVD idea to iterate—using the K-SVD output to re-run NLM and re-train the dictionary—could further improve results. Applying the combined pipeline to colour images or medical image reconstruction are further promising directions.

REFERENCES

- [1] M. Aharon, M. Elad, and A. Bruckstein, “K-SVD: An algorithm for designing overcomplete dictionaries for sparse representation,” *IEEE Trans. Signal Process.*, vol. 54, no. 11, pp. 4311–4322, Nov. 2006.
- [2] K. Engan, S. O. Aase, and J. H. Husoy, “Method of optimal directions for frame design,” in *Proc. IEEE ICASSP*, 1999, pp. 2443–2446.
- [3] Y. C. Pati, R. Rezaifar, and P. S. Krishnaprasad, “Orthogonal matching pursuit: Recursive function approximation with applications to wavelet decomposition,” in *Proc. 27th Asilomar Conf. Signals, Systems Comput.*, 1993, pp. 40–44.
- [4] S. G. Mallat and Z. Zhang, “Matching pursuits with time-frequency dictionaries,” *IEEE Trans. Signal Process.*, vol. 41, no. 12, pp. 3397–3415, Dec. 1993.
- [5] B. A. Olshausen and D. J. Field, “Emergence of simple-cell receptive field properties by learning a sparse code for natural images,” *Nature*, vol. 381, pp. 607–609, 1996.
- [6] M. Elad and M. Aharon, “Image denoising via sparse and redundant representations over learned dictionaries,” *IEEE Trans. Image Process.*, vol. 15, no. 12, pp. 3736–3745, Dec. 2006.
- [7] F. Pedregosa *et al.*, “Scikit-learn: Machine learning in Python,” *J. Mach. Learn. Res.*, vol. 12, pp. 2825–2830, 2011.
- [8] A. Buades, B. Coll, and J.-M. Morel, “A non-local algorithm for image denoising,” in *Proc. IEEE CVPR*, 2005, vol. 2, pp. 60–65.